

VISVESVARAYA TECHNOLOGICAL UNIVERSITY

“JnanaSangama”, Belgaum -590014, Karnataka.



LAB REPORT on

Machine Learning (23CS6PCMAL)

Submitted by

SAKSHI SHETTY (1BM22CS234)

in partial fulfillment for the award of the degree of

BACHELOR OF ENGINEERING
in
COMPUTER SCIENCE AND ENGINEERING



B.M.S. COLLEGE OF ENGINEERING

(Autonomous Institution under VTU)

BENGALURU-560019

Sep-2024 to Jan-2025

B.M.S. College of Engineering,
Bull Temple Road, Bangalore 560019
(Affiliated To Visvesvaraya Technological University, Belgaum)

Department of Computer Science and Engineering



CERTIFICATE

This is to certify that the Lab work entitled “Machine Learning (23CS6PCMAL)” carried out by **Sakshi Shetty (1BM22CS234)**, who is a bonafide student of **B.M.S. College of Engineering**. It is in partial fulfillment for the award of **Bachelor of Engineering in Computer Science and Engineering** of the Visvesvaraya Technological University, Belgaum. The Lab report has been approved as it satisfies the academic requirements in respect of an Machine Learning (23CS6PCMAL) work prescribed for the said degree.

Lab Faculty Incharge Name: Ms. Saritha A N Assistant Professor Department of CSE, BMSCE	Dr. Kavitha Sooda Professor & HOD Department of CSE, BMSCE
--	---

Index

Sl. No.	Date	Experiment Title	Page No.
1	5-3-2025	Write a python program to import and export data using Pandas library functions	
2	5-3-2025	Demonstrate various data pre-processing techniques for a given dataset	
3	12-3-2025	Implement Linear and Multi-Linear Regression algorithm using appropriate dataset	
4	19-3-2025	Build Logistic Regression Model for a given dataset	
5	27-3-2025	Use an appropriate data set for building the decision tree (ID3) and apply this knowledge to classify a new sample	
6	7-4-2025	Build KNN Classification model for a given dataset	
7	21-4-2025	Build Support vector machine model for a given dataset	
8	16-4-2025	Implement Random forest ensemble method on a given dataset	
9	7-5-2025	Implement Boosting ensemble method on a given dataset	
10	7-5-2025	Build k-Means algorithm to cluster a set of data stored in a .CSV file	
11	12-5-2025	Implement Dimensionality reduction using Principal Component Analysis (PCA) method	

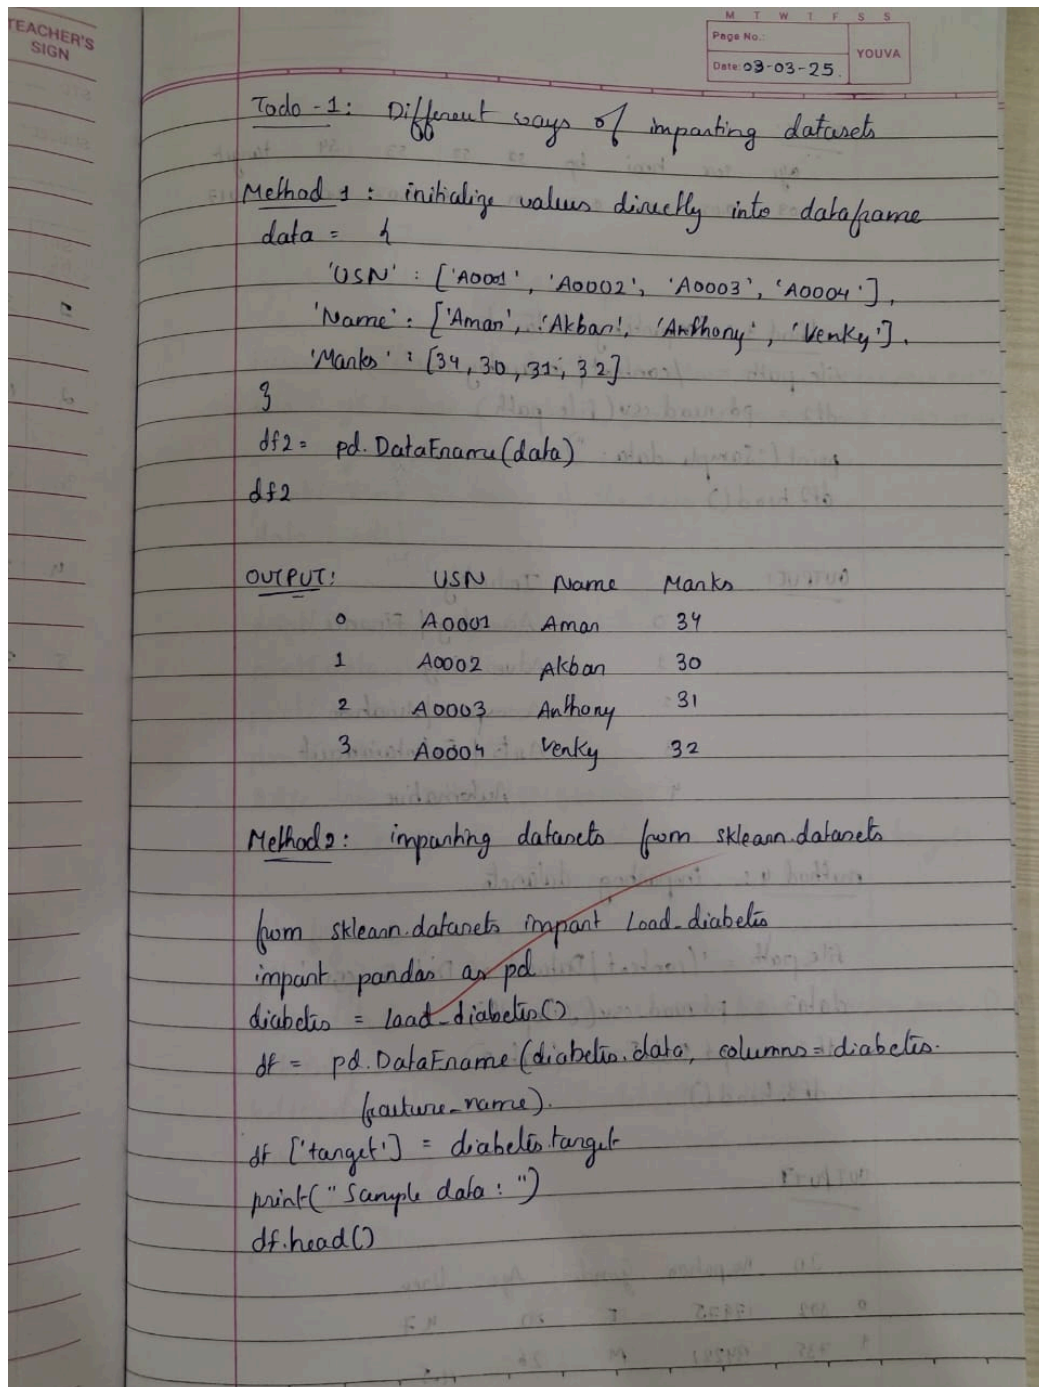
Github Link:

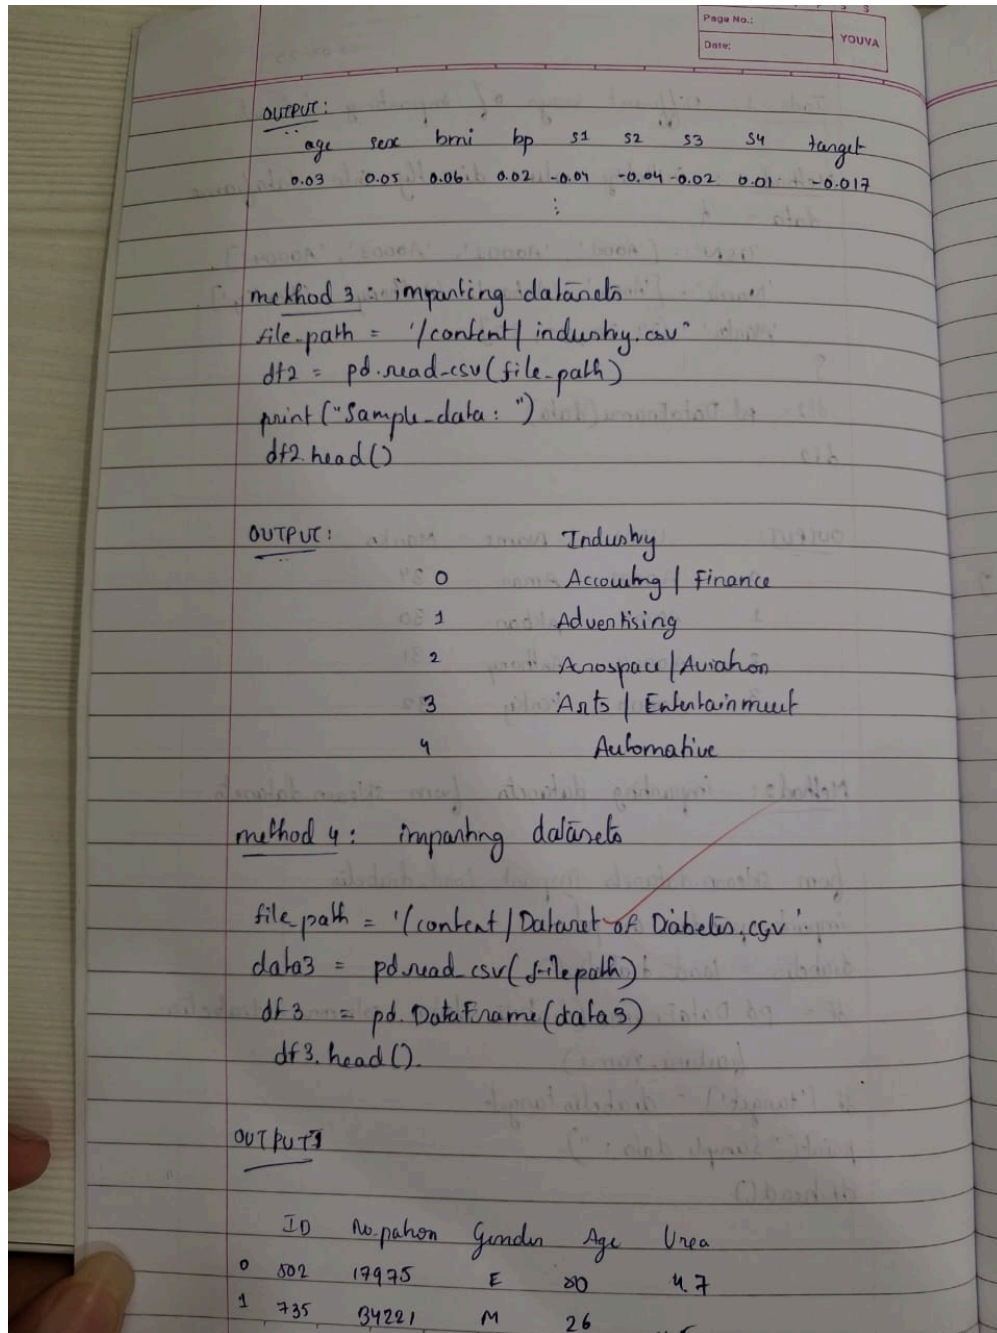
<https://github.com/Sakshishetty24/ML-1BM22CS234>

Program 1

Write a python program to import and export data using Pandas library functions

Screenshot





Code:

```
from sklearn.datasets import load_iris
import pandas as pd

iris = load_iris()
df = pd.DataFrame(iris.data, columns=iris.feature_names)
df['target'] = iris.target
print("sample data: ")
df.head()
```

```

# method 1
data = {
    'USN': ['A001', 'A002', 'A003', 'A004'],
    'Name': ['Amar', 'Akbar', 'Anthony', 'Venky'],
    'Marks': [34, 30, 31, 32]
}

df2 = pd.DataFrame(data)
df2

# method 2
from sklearn.datasets import load_diabetes
import pandas as pd

diabetes = load_diabetes()
df = pd.DataFrame(diabetes.data, columns=diabetes.feature_names)
df['target'] = diabetes.target
print("sample data: ")
df.head()

# method 3

# Load data from a CSV file (replace 'data.csv' with your file path)
file_path = '/content/industry.csv' # Ensure the file exists in the same
directory
df2 = pd.read_csv(file_path)
print("Sample data:")
df2.head()

# method 4

file_path = '/content/Dataset of Diabetes .csv'
data3 = pd.read_csv(file_path)
df3 = pd.DataFrame(data3)

df3

#Using the code given in the above slides, do the exercise of the "Stock
Market Data Analysis", considering the follwoing
# 1. HDFC Bank Ltd. , ICICI Bank Ltd , Kotak Mahindra Bank Ltd.
# tickers = ["HDFCBANK.NS", "ICICIBANK.NS", "KOTAKBANK.NS"]

```

```

# 2. Start date: 2024-01-01, End date: 2024-12-30

import yfinance as yf
import pandas as pd
import matplotlib.pyplot as plt

tickers = ["HDFCBANK.NS", "ICICIBANK.NS", "KOTAKBANK.NS"]
data = yf.download(tickers, start="2024-01-01", end="2024-12-30",
group_by='ticker')
print("First 5 rows of the dataset:")
data.head()

print("\nShape of the dataset:")
print(data.shape)
print("\nColumn names:")
print(data.columns)
hdfc_data = data['HDFCBANK.NS']
print("\nSummary statistics for Reliance Industries:")
print(hdfc_data.describe())
hdfc_data['Daily Return'] = hdfc_data['Close'].pct_change()

# icici bank
icici_data = data['ICICIBANK.NS']
print(hdfc_data.describe())
icici_data['Daily Return'] = icici_data['Close'].pct_change()

# Kotak bank
kotak_data = data['KOTAKBANK.NS']
print(hdfc_data.describe())
kotak_data['Daily Return'] = kotak_data['Close'].pct_change()

plt.figure(figsize=(12, 6))
plt.subplot(2, 1, 1)
hdfc_data['Close'].plot(title="HDFC bank - Closing Price")
plt.subplot(2, 1, 2)
hdfc_data['Daily Return'].plot(title="HDFC bank - Daily Returns",
color='orange')
plt.tight_layout()
plt.show()

plt.figure(figsize=(12, 6))
plt.subplot(2, 1, 1)

```

```
icici_data['Close'].plot(title="ICICI bank - Closing Price")
plt.subplot(2, 1, 2)
icici_data['Daily Return'].plot(title="ICICI bank - Daily Returns",
color='orange')
plt.tight_layout()
plt.show()

plt.figure(figsize=(12, 6))
plt.subplot(2, 1, 1)
kotak_data['Close'].plot(title="KOTAK bank - Closing Price")
plt.subplot(2, 1, 2)
kotak_data['Daily Return'].plot(title="KOTAK bank - Daily Returns",
color='orange')
plt.tight_layout()
plt.show()

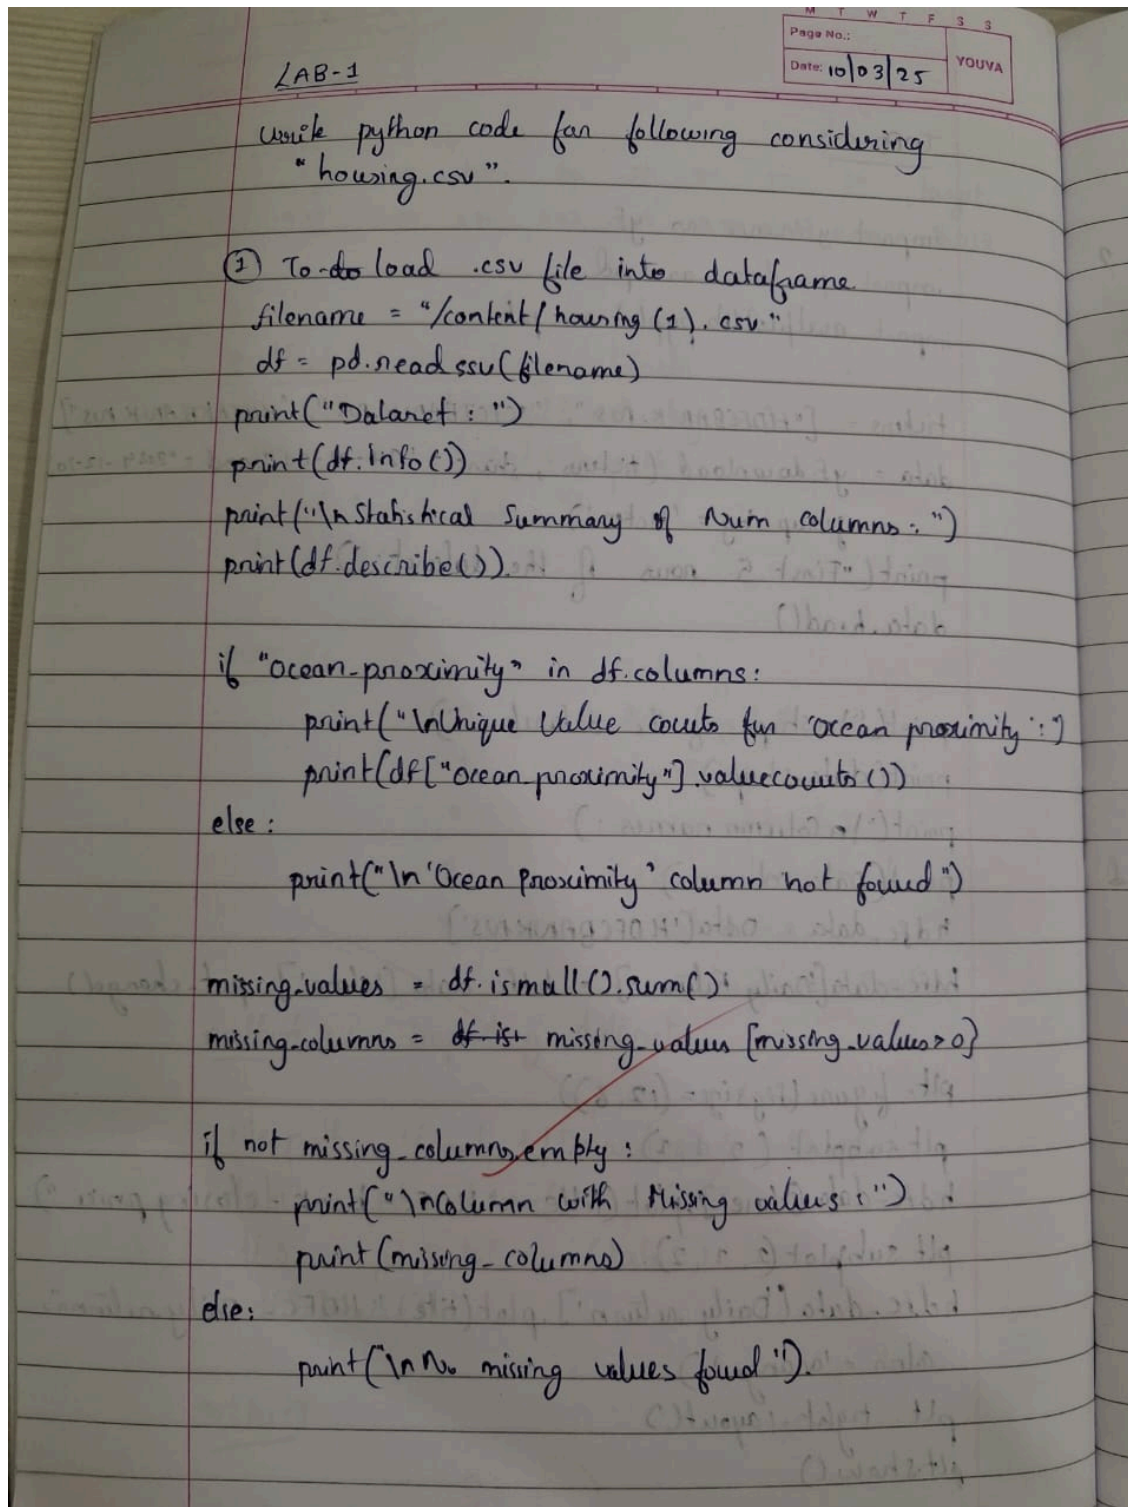
# Save the Reliance data to a CSV file

hdfc_data.to_csv('hdfc_stock_data.csv')
icici_data.to_csv('icici_stock_data.csv')
kotak_data.to_csv('kotak_stock_data.csv')
```


Program 2

Demonstrate various data pre-processing techniques for a given dataset

Screenshot



Questions:

Q1) Which columns in the datasets had missing values?
How did you handle them?

Soln: missing value columns:

Adult income dataset \rightarrow Age, salary

Diabetes dataset \rightarrow Glucose, BMI

Handling approach:

Adult income ds \rightarrow Fange - used median since it

less sensitive to outliers

\rightarrow For salary - used mean as

salaries typically follow normal distributions

diabetes dataset \rightarrow Glucose - used median since

glucose levels may have outliers

BMI - used mean assuming normal distributions

Q2) Which categorical columns did you identify in the dataset? How did you encode them?

Soln:

Adult income Dataset:

categorical columns - Gender \rightarrow original encoding

city \rightarrow one-hot encoding

Diabetes dataset

categorical columns - Gender \rightarrow original encoding

Outcome \rightarrow one-hot encoding

Q3) what is the difference b/w min-max scaling and standardization? when would you use one over the other?

Soln

Min-max scaling

$$\rightarrow x' = \frac{x - x_{\min}}{x_{\max} - x_{\min}}$$

→ Scales values b/w 0 & 1

→ Sensitive to outliers

Standardization

$$\rightarrow x' = \frac{x - \mu}{\sigma}$$

→ Transforms data to have mean = 0 & variance = 1

→ Less affected by outliers

→ when data is not normally distributed and has known bounds, min-max scaling is used

→ when data follows a normal distribution, Standardization is used

Shreya B
 10/3/25

code:

```
# -*- coding: utf-8 -*-
"""LAB-1.ipynb

Automatically generated by Colab.

Original file is located at
    https://colab.research.google.com/drive/1LFiPSjr6wkzvYXycyOlrEerHW0HtTt12
"""

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import OrdinalEncoder, OneHotEncoder
from sklearn.preprocessing import StandardScaler, MinMaxScaler
from scipy import stats

def createdata():
    data = {
        'Age': np.random.randint(18, 70, size=20),
        'Salary': np.random.randint(30000, 120000, size=20),
        'Purchased': np.random.choice([0, 1], size=20),
        'Gender': np.random.choice(['Male', 'Female'], size=20),
        'City': np.random.choice(['New York', 'San Francisco', 'Los Angeles'],
size=20)
    }

    df = pd.DataFrame(data)
    return df

df = createdata()
df.head(10)

df.shape

# Introduce some missing values for demonstration
df.loc[5, 'Age'] = np.nan
df.loc[10, 'Salary'] = np.nan
df.head(10)
```

```

# Basic information about the dataset
print(df.info())

# Summary statistics
print(df.describe())

#Code to Find Missing Values
# Check for missing values in each column
missing_values = df.isnull().sum()

# Display columns with missing values
print(missing_values[missing_values > 0])

#Set the values to some value (zero, the mean, the median, etc.).
# Step 1: Create an instance of SimpleImputer with the median strategy for
Age and mean strategy for Salary
imputer1 = SimpleImputer(strategy="median")
imputer2 = SimpleImputer(strategy="mean")

df_copy=df

# Step 2: Fit the imputer on the "Age" and "Salary" column
# Note: SimpleImputer expects a 2D array, so we reshape the column
imputer1.fit(df_copy[["Age"]])
imputer2.fit(df_copy[["Salary"]])

# Step 3: Transform (fill) the missing values in the "Age" and "Salary" column
df_copy["Age"] = imputer1.transform(df[["Age"]])
df_copy["Salary"] = imputer2.transform(df[["Salary"]])

# Verify that there are no missing values left
print(df_copy["Age"].isnull().sum())
print(df_copy["Salary"].isnull().sum())

#Handling Categorical Attributes
#Using Ordinal Encoding for gender Column and One-Hot Encoding for City
Column# Initialize OrdinalEncoder
ordinal_encoder = OrdinalEncoder(categories=[["Male", "Female"]])
# Fit and transform the data

```



```

df_copy["Gender_Encoded"] =
ordinal_encoder.fit_transform(df_copy[["Gender"]])

# Initialize OneHotEncoder
onehot_encoder = OneHotEncoder()

# Fit and transform the "City" column
encoded_data = onehot_encoder.fit_transform(df[["City"]])

# Convert the sparse matrix to a dense array
encoded_array = encoded_data.toarray()

# Convert to DataFrame for better visualization
encoded_df = pd.DataFrame(encoded_array,
columns=onehot_encoder.get_feature_names_out(["City"]))
df_encoded = pd.concat([df_copy, encoded_df], axis=1)

df_encoded.drop("Gender", axis=1, inplace=True)
df_encoded.drop("City", axis=1, inplace=True)

df_encoded.head()

#Data Transformation
# Min-Max Scaler/Normalization (range 0-1)
#Pros: Keeps all data between 0 and 1; ideal for distance-based models.
#Cons: Can distort data distribution, especially with extreme outliers.
normalizer = MinMaxScaler()
df_encoded[['Salary']] = normalizer.fit_transform(df_encoded[['Salary']])
df_encoded.head()

# Standardization (mean=0, variance=1)
#Pros: Works well for normally distributed data; suitable for many models.
#Cons: Sensitive to outliers.
scaler = StandardScaler()
df_encoded[['Age']] = scaler.fit_transform(df_encoded[['Age']])
df_encoded.head()

#Removing Outliers
# Outlier Detection and Treatment using IQR.
#Pros: Simple and effective for mild outliers.
#Cons: May overly reduce variation if there are many extreme outliers.

```

```

df_encoded_copy1=df_encoded
df_encoded_copy2=df_encoded
df_encoded_copy3=df_encoded

Q1 = df_encoded_copy1['Salary'].quantile(0.25)
Q3 = df_encoded_copy1['Salary'].quantile(0.75)
IQR = Q3 - Q1
lower_bound = Q1 - 1.5 * IQR
upper_bound = Q3 + 1.5 * IQR
df_encoded_copy1['Salary'] = np.where(df_encoded_copy1['Salary'] >
upper_bound, upper_bound,
                                     np.where(df_encoded_copy1['Salary'] < lower_bound,
lower_bound, df_encoded_copy1['Salary']))

df_encoded_copy1.head()

#Removing Outliers
# Z-score method
#Pros: Good for normally distributed data.
#Cons: Not suitable for non-normal data; may miss outliers in skewed
distributions.

df_encoded_copy2['Salary_zscore'] = stats.zscore(df_encoded_copy2['Salary'])
df_encoded_copy2['Salary'] = np.where(df_encoded_copy2['Salary_zscore'].abs()
> 3, np.nan, df_encoded_copy2['Salary']) # Replace outliers with NaN
df_encoded_copy2.head()

#Removing Outliers
# Median replacement for outliers
#Pros: Keeps distribution shape intact, useful when capping isn't feasible.
#Cons: May distort data if outliers represent real phenomena.
df_encoded_copy3['Salary_zscore'] = stats.zscore(df_encoded_copy3['Salary'])
median_salary = df_encoded_copy3['Salary'].median()
df_encoded_copy3['Salary'] = np.where(df_encoded_copy3['Salary_zscore'].abs()
> 3, median_salary, df_encoded_copy3['Salary'])
df_encoded_copy3.head()

'''
At the start of the Lab, in the Observation book, Write python code for the
following considering filename as "housing.csv"
i. To load .csv file into the data frame

```

```

ii. To display information of all columns
iii. To display statistical information of all numerical
iv. To display the count of unique labels for "Ocean Proximity" column
v. To display which attributes (columns) in a dataset have missing values
count greater than zero
Step-2: Show the observation book to lab batch faculty incharge.
Step-3: Do the "To Do" tasks given in the PPT
Step-4: At the end of the lab,
i. Write the answers for questions given in the PPT and show it to lab batch
faculty incharge
ii. Should upload the code in your respective GitHub account.
File name format:yourUSN_Lab-1-DataProcessing.ipynb
'''

filename = "/content/housing (1).csv"
df = pd.read_csv(filename)

print("Dataset Information:")
print(df.info())

print("\nStatistical Summary of Numerical Columns:")
print(df.describe())

if "ocean_proximity" in df.columns:
    print("\nUnique Value Counts for 'Ocean Proximity':")
    print(df["ocean_proximity"].value_counts())
else:
    print("\n'Ocean Proximity' column not found in the dataset.")

missing_values = df.isnull().sum()
missing_columns = missing_values[missing_values > 0]

if not missing_columns.empty:
    print("\nColumns with Missing Values:")
    print(missing_columns)
else:
    print("\nNo missing values found in the dataset.")

data2 = pd.read_csv("/content/Dataset_with_Nulls.csv")
data2.head()

```



```
data2.info()

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import OrdinalEncoder, OneHotEncoder,
StandardScaler, MinMaxScaler
from scipy import stats

# Load dataset
file_path = "/content/Dataset_with_Nulls.csv"
df = pd.read_csv(file_path)

### Step 1: Handling Missing Values ###
# Identify missing values
print("Missing values before handling:\n", df.isnull().sum())

# Handling missing numerical columns - Median for 'Age', Mean for other
numerical values
num_imputer_median = SimpleImputer(strategy="median")
num_imputer_mean = SimpleImputer(strategy="mean")

df["AGE"] = num_imputer_median.fit_transform(df[["AGE"]])
for col in ["Urea", "Cr", "HbA1c", "Chol", "TG", "HDL", "LDL", "VLDL",
"BMI"]:
    df[col] = num_imputer_mean.fit_transform(df[[col]])

# Convert categorical columns to string type
df["Gender"] = df["Gender"].astype(str)
df["CLASS"] = df["CLASS"].astype(str)

# Handling missing categorical columns - Fill with Mode
cat_imputer = SimpleImputer(strategy="most_frequent")

df["Gender"] = cat_imputer.fit_transform(df[["Gender"]]).ravel()
df["CLASS"] = cat_imputer.fit_transform(df[["CLASS"]]).ravel()

print("Missing values after handling:\n", df.isnull().sum())
```

```

### Step 2: Handling Categorical Attributes ###
# Encode 'Gender' using Ordinal Encoding (Male = 0, Female = 1)
ordinal_encoder = OrdinalEncoder(categories=[["Male", "Female"]])
df['Gender'] = df['Gender'].replace({'F': 'Female', 'f': 'Female', 'M': 'Male'}) # Handle variations
# Replace NaN with Mode (most frequent) for Gender
df['Gender'] = df['Gender'].fillna(df['Gender'].mode()[0]) # Fill na with mode if any

# Now apply ordinal encoding
ordinal_encoder = OrdinalEncoder(categories=[["Male", "Female"]])
# One-Hot Encoding for 'CLASS'
onehot_encoder = OneHotEncoder(sparse_output=False)
class_encoded = onehot_encoder.fit_transform(df[["CLASS"]])

# Convert to DataFrame
class_encoded_df = pd.DataFrame(class_encoded,
                                columns=onehot_encoder.get_feature_names_out(["CLASS"]))

# Merge One-Hot Encoded Data and drop original categorical columns
df = pd.concat([df, class_encoded_df], axis=1)
df.drop(["Gender", "CLASS"], axis=1, inplace=True)

### Step 3: Data Transformation ###
# Min-Max Scaling for Salary
minmax_scaler = MinMaxScaler()
df[["Urea", "Cr", "HbA1c", "Chol", "TG", "HDL", "LDL", "VLDL", "BMI"]] = minmax_scaler.fit_transform(
    df[["Urea", "Cr", "HbA1c", "Chol", "TG", "HDL", "LDL", "VLDL", "BMI"]]
)

# Standardization for Age
standard_scaler = StandardScaler()
df[["AGE"]] = standard_scaler.fit_transform(df[["AGE"]])

### Step 4: Removing Outliers ###
# IQR Method for 'Salary' (Replacing Outliers with Boundaries)
for col in ["Urea", "Cr", "HbA1c", "Chol", "TG", "HDL", "LDL", "VLDL", "BMI"]:
    Q1 = df[col].quantile(0.25)
    Q3 = df[col].quantile(0.75)

```

```
IQR = Q3 - Q1
lower_bound = Q1 - 1.5 * IQR
upper_bound = Q3 + 1.5 * IQR
df[col] = np.where(df[col] > upper_bound, upper_bound, np.where(df[col] <
lower_bound, lower_bound, df[col]))

# Z-score method for 'AGE' (Replacing Outliers with NaN)
df["AGE_zscore"] = stats.zscore(df["AGE"])
df["AGE"] = np.where(df["AGE_zscore"].abs() > 3, np.nan, df["AGE"])

# Median Replacement for Outliers in 'AGE'
median_age = df["AGE"].median()
df["AGE"] = np.where(df["AGE"].isnull(), median_age, df["AGE"])

# Drop auxiliary columns
df.drop(columns=["AGE_zscore"], inplace=True)

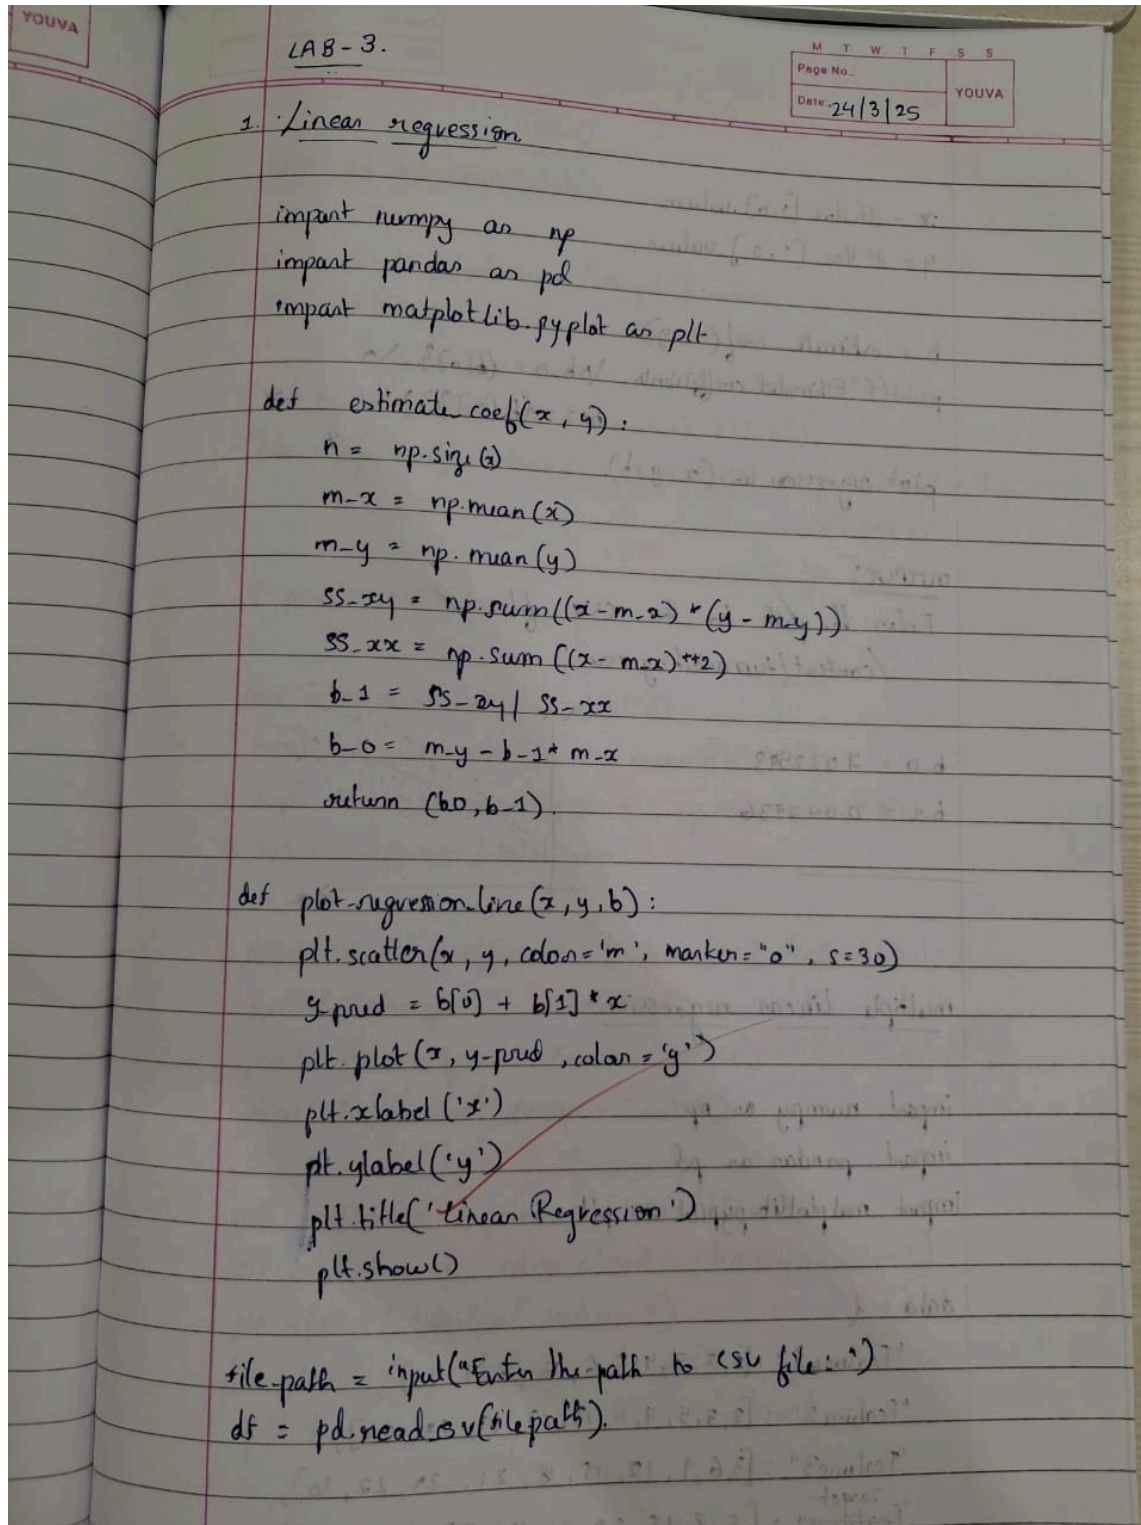
df

print("Preprocessing Complete. Cleaned dataset saved!")
```

Program 3

Implement Linear and Multi-Linear Regression algorithm using appropriate dataset

Screenshot



```
x = df.iloc[:,0].values
y = df.iloc[:,1].values
```

```
b = estimate_coef(x, y)
print(f"Estimated coefficients: \nb_0 = {b[0]} \nb_1 = {b[1]}")
```


```
plot_regression_line(x, y, b)
```

OUTPUT:

Enter the path to the csv file
(content/tumanking.csv)

$b_0 = 7.032593$

$b_1 = 0.047536$



multiple linear regression:

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
```

```
data = {
    "Feature 1": [1, 2, 3, 4, 5, 6, 7, 8, 9, 10],
    "Feature 2": [2, 3, 5, 7, 11, 13, 17, 19, 23, 29],
    "Feature 3": [3, 6, 9, 12, 15, 18, 21, 24, 27, 30],
    "Target": [5, 9, 15, 22, 31, 41, 53, 66, 80, 96],
}
```

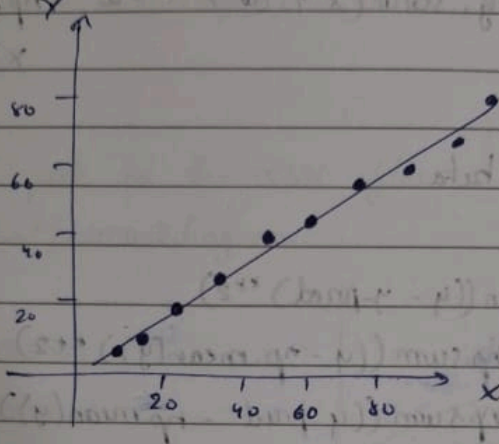

output:

Model coefficients: [0.402 3.313 0.120]

Intercept: -9.159950

MSE: 5.36291281

R-square: 0.99353261



3. logistic regression

```
import numpy as np
```

```
import matplotlib.pyplot as plt
```

```
def sigmoid(z):
```

```
    return 1 / (1 + np.exp(-z))
```

```
def compute_cost(X, y, theta):
```

```
    m = len(y)
```

```
    h = sigmoid(X @ theta)
```

```
    cost = (-1/m) * np.sum(y * np.log(h) + (1 - y) *  
        np.log(1 - h))
```

```
    return cost
```

code:

```
import numpy as np
import matplotlib.pyplot as plt

def estimate_coef(x, y):
    # number of observations/points
    n = np.size(x)

    # mean of x and y vector
    m_x = np.mean(x)
    m_y = np.mean(y)

    # calculating cross-deviation and deviation about x
    SS_xy = np.sum((x - m_x) * (y - m_y))
    SS_xx = np.sum((x - m_x) ** 2)

    # calculating regression coefficients
    b_1 = SS_xy / SS_xx
    b_0 = m_y - b_1 * m_x

    return (b_0, b_1)

def plot_regression_line(x, y, b):
    # plotting the actual points as scatter plot
    plt.scatter(x, y, color="m", marker="o", s=30)

    # predicted response vector
    y_pred = b[0] + b[1] * x

    # plotting the regression line
    plt.plot(x, y_pred, color="g")

    # putting labels
    plt.xlabel('x')
    plt.ylabel('y')
    plt.title("Linear Regression")
    plt.show()

x = np.array([0, 1, 2, 3, 4, 5, 6, 7, 8, 9])
y = np.array([1, 3, 2, 5, 7, 8, 8, 9, 10, 12])
```

```

# estimating coefficients
b = estimate_coef(x, y)
print(f"Estimated coefficients:\nb_0 = {b[0]} \nb_1 = {b[1]}")

# plot regression line
plot_regression_line(x, y, b)

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

def estimate_coef(x, y):
    n = np.size(x)
    m_x = np.mean(x)
    m_y = np.mean(y)
    SS_xy = np.sum((x - m_x) * (y - m_y))
    SS_xx = np.sum((x - m_x) ** 2)
    b_1 = SS_xy / SS_xx
    b_0 = m_y - b_1 * m_x
    return (b_0, b_1)

def plot_regression_line(x, y, b):
    plt.scatter(x, y, color="m", marker="o", s=30)
    y_pred = b[0] + b[1] * x
    plt.plot(x, y_pred, color="g")
    plt.xlabel('x')
    plt.ylabel('y')
    plt.title("Linear Regression")
    plt.show()

# Load dataset
file_path = input("Enter the path to the CSV file: ")
df = pd.read_csv(file_path)

# Assuming the dataset has two numerical columns: 'x' and 'y'
x = df.iloc[:, 0].values # First column as x
y = df.iloc[:, 1].values # Second column as y

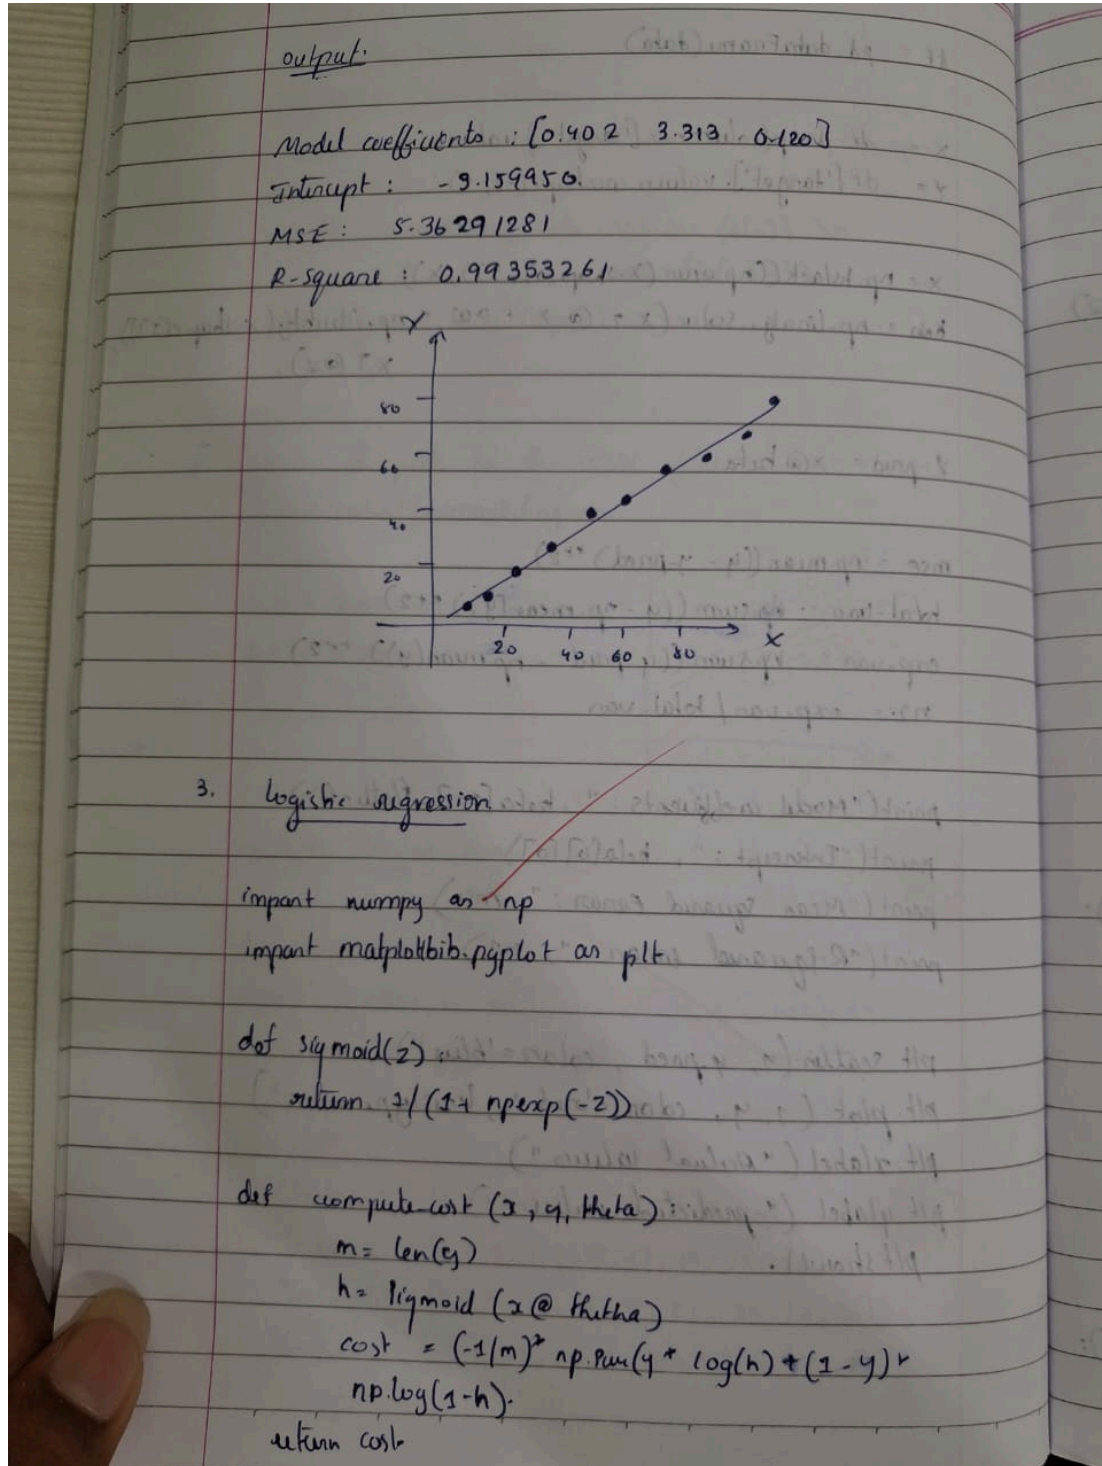
b = estimate_coef(x, y)
print(f"Estimated coefficients:\nb_0 = {b[0]} \nb_1 = {b[1]}")
plot_regression_line(x, y, b)

```


Program 4

Build Logistic Regression Model for a given dataset

Screenshot



code:

```
import numpy as np
import matplotlib.pyplot as plt

def sigmoid(z):
    return 1 / (1 + np.exp(-z))

def compute_cost(X, y, theta):
    m = len(y)
    h = sigmoid(X @ theta)
    cost = (-1/m) * np.sum(y * np.log(h) + (1 - y) * np.log(1 - h))
    return cost

def gradient_descent(X, y, theta, alpha, iterations):
    m = len(y)
    cost_history = []

    for _ in range(iterations):
        gradient = (1/m) * X.T @ (sigmoid(X @ theta) - y)
        theta -= alpha * gradient
        cost_history.append(compute_cost(X, y, theta))

    return theta, cost_history

def predict(X, theta):
    return (sigmoid(X @ theta) >= 0.5).astype(int)

# Generate synthetic binary classification data
np.random.seed(42)
X = np.random.rand(100, 1) * 10 # Feature values between 0 and 10
y = (X > 5).astype(int).ravel() # Label: 1 if X > 5, else 0

# Add intercept term
X_b = np.c_[np.ones((X.shape[0], 1)), X]

# Initialize parameters
theta = np.zeros(X_b.shape[1])
alpha = 0.1
iterations = 1000

# Train logistic regression using gradient descent
```

```
theta, cost_history = gradient_descent(X_b, y, theta, alpha, iterations)

# Make predictions
y_pred = predict(X_b, theta)

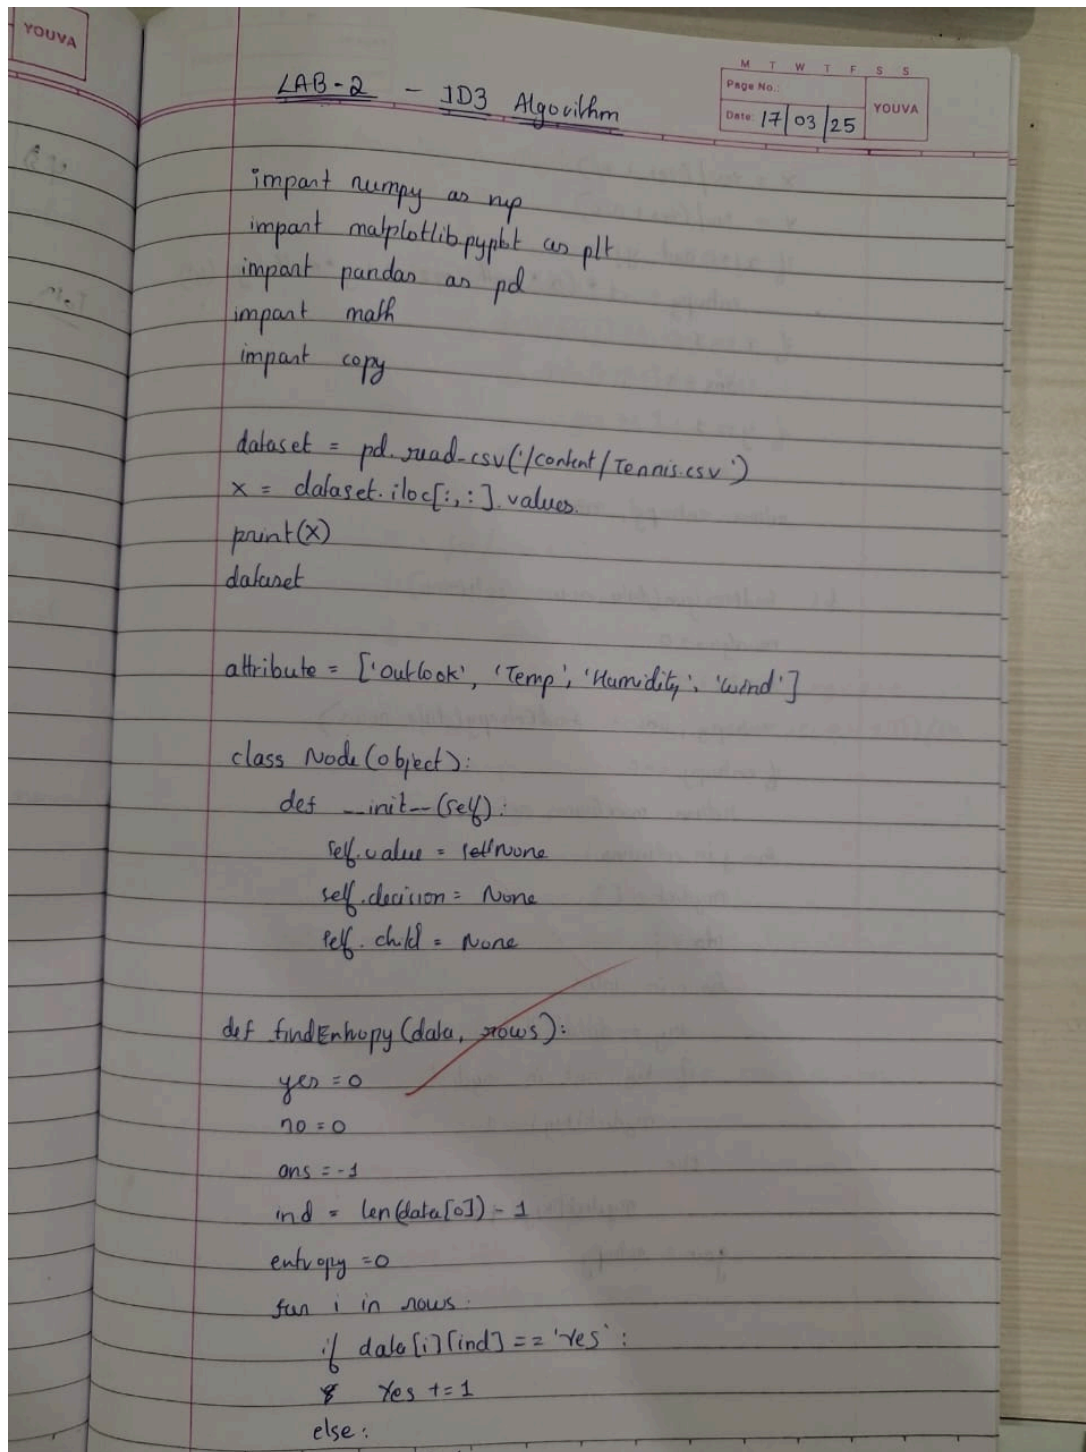
# Compute accuracy
accuracy = np.mean(y_pred == y)
print(f"Accuracy: {accuracy:.2f}")

# Plot the decision boundary
plt.scatter(X, y, color='blue', label='Actual Data')
plt.scatter(X, y_pred, color='red', marker='x', label='Predicted Labels')
plt.xlabel("Feature X")
plt.ylabel("Class (0 or 1)")
plt.legend()
plt.title("Logistic Regression Model (Without Scikit-learn)")
plt.show()
```

Program 5

Use an appropriate data set for building the decision tree (ID3) and apply this knowledge to classify a new sample

Screenshot



The image shows a handwritten code snippet in a notebook. The notebook has a header with 'YOUVA' on the left and a date stamp '17/03/25' on the right. The title 'LAB-2 - ID3 Algorithm' is written at the top. The code is written in Python and includes imports for numpy, matplotlib.pyplot, pandas, math, and copy. It reads a CSV file 'Tennis.csv' and prints its values. It also defines a 'Node' class and a 'findEntropy' function. The 'findEntropy' function is partially visible, showing the initialization of 'yes' and 'no' counts and the calculation of the index 'ind'.

```
LAB-2 - ID3 Algorithm

import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
import math
import copy

dataset = pd.read_csv('/content/Tennis.csv')
x = dataset.iloc[:, :].values
print(x)
dataset

attribute = ['outlook', 'Temp', 'Humidity', 'wind']

class Node(object):
    def __init__(self):
        self.value = self.none
        self.decision = None
        self.child = None

def findEntropy(data, rows):
    yes = 0
    no = 0
    ans = -1
    ind = len(data[0]) - 1
    entropy = 0
    for i in rows:
        if data[i][ind] == 'yes':
            yes += 1
        else:
```

$g_2(y)$

```

for key in mydict:
    yes = 0
    no = 0
    for k in rows:
        if data[k][j] == key:
            if data[k][j-1] == 'yes':
                yes += 1
            else:
                no += 1
    x = yes / (yes + no)
    y = no / (yes + no)
    if x != 0 and y != 0:
        gain += (mydict[key] * (x * math.log2(x) +
            y * math.log2(y))) / 14

    if gain > maxgain:
        maxgain = gain
        newidx = j

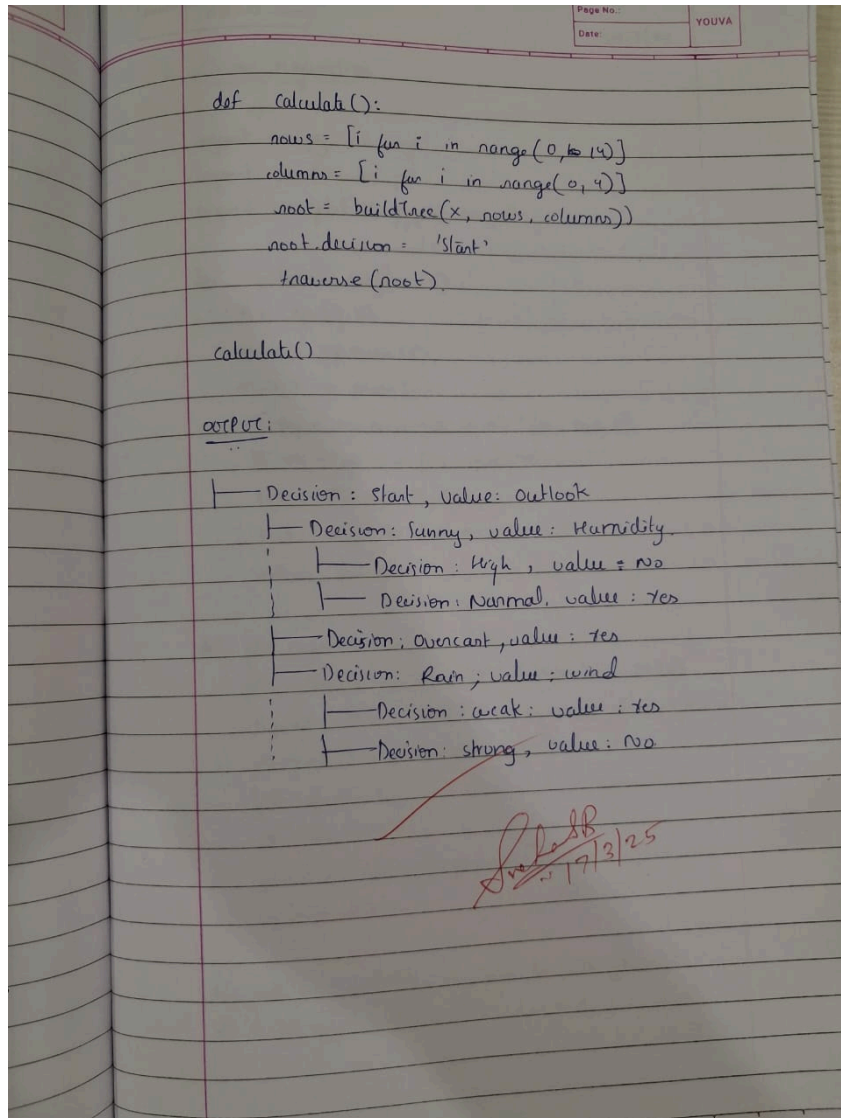
return maxgain, newidx, ans

def buildTree(data, rows, columns):
    maxgain, idx, ans = findMaxGain(x, rows, columns)
    root = node()
    root.children = []
    if maxgain == 0:
        if ans == 1:
            root.value = 'yes'
        else:
            root.value = 'no'
    return root
  
```

```

x = yes / (yes + no)
y = no / (yes + no)
if x != 0 and y != 0:
    entropy = -1 * (x * math.log2(x) + y * math.log2(y))
    if x == 1:
        ans = 1
    if y == 1:
        ans = 0
    return entropy, ans

def findMaxGain(data, rows, columns):
    maxgain = 0
    newidx = -1
    entropy, ans = findEntropy(data, rows)
    if entropy == 0:
        return maxgain, newidx, ans
    for j in columns:
        mydict = {}
        idx = j
        for i in rows:
            key = data[i][idx]
            if key not in mydict:
                mydict[key] = 1
            else:
                mydict[key] += 1
        gain = entropy
  
```



code:

```
import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
import math
import copy

dataset = pd.read_csv('/content/Tennis.csv')
X = dataset.iloc[:, :].values
print(X)

dataset

attribute = ['Outlook', 'Temp', 'Humidity', 'Wind']
```



```

class Node(object):
    def __init__(self):
        self.value = None
        self.decision = None
        self.child = None

def findEntropy(data, rows):
    yes=0
    no=0
    ans=-1
    idx=len(data[0])-1
    entropy=0

    for i in rows:
        if data[i][idx]=='Yes':
            yes=yes+1
        else:
            no=no+1

    x=yes/(yes+no)
    y=no/(yes+no)
    if x!=0 and y!=0:
        entropy= -1*(x*math.log2(x)+y*math.log2(y))
    if x==1:
        ans = 1
    if y==1:
        ans = 0
    return entropy, ans

def findMaxGain(data, rows, columns):
    maxGain = 0
    retidx = -1
    entropy, ans = findEntropy(data, rows)
    if entropy == 0:
        """if ans == 1:
            print("Yes")
        else:
            print("No")"""
    return maxGain, retidx, ans
    for j in columns:

```

```

mydict = {}
idx = j
for i in rows:
    key = data[i][idx]
    if key not in mydict:
        mydict[key] = 1
    else:
        mydict[key] = mydict[key] + 1
gain = entropy

# print(mydict)
for key in mydict:
    yes = 0
    no = 0
    for k in rows:
        if data[k][j] == key:
            if data[k][-1] == 'Yes':
                yes = yes + 1
            else:
                no = no + 1
    # print(yes, no)
    x = yes/(yes+no)
    y = no/(yes+no)
    # print(x, y)
    if x != 0 and y != 0:
        gain += (mydict[key] * (x*math.log2(x) + y*math.log2(y)))/14
# print(gain)
if gain > maxGain:
    # print("hello")
    maxGain = gain
    retidx = j

return maxGain, retidx, ans

def buildTree(data, rows, columns):

    maxGain, idx, ans = findMaxGain(X, rows, columns)
    root = Node()
    root.childs = []
    # print(maxGain)

```



```

    if maxGain == 0:
        if ans == 1:
            root.value = 'Yes'
        else:
            root.value = 'No'
        return root

    root.value = attribute[idx]
    mydict = {}
    for i in rows:
        key = data[i][idx]
        if key not in mydict:
            mydict[key] = 1
        else:
            mydict[key] += 1

    newcolumns = copy.deepcopy(columns)
    newcolumns.remove(idx)
    for key in mydict:
        newrows = []
        for i in rows:
            if data[i][idx] == key:
                newrows.append(i)
        # print(newrows)
        temp = buildTree(data, newrows, newcolumns)
        temp.decision = key
        root.chlds.append(temp)
    return root

def traverse(root, level=0):
    indent = "    " * level
    print(f"{indent}|— Decision: {root.decision}, Value: {root.value}")

    for i, child in enumerate(root.chlds):
        # if i == len(root.chlds) - 1:
        #     traverse(child, level + 1)
        # else:
        traverse(child, level + 1)

def calculate():
    rows = [i for i in range(0, 14)]

```

```

        columns = [i for i in range(0, 4)]
        root = buildTree(X, rows, columns)
        root.decision = 'Start'
        traverse(root)

calculate()

from graphviz import Source

dot_code = """
digraph G {
    edge [dir=forward]
    node [shape=box, style=bold]

    A [label="OUTLOOK"]
    B [label="HUMIDITY"]
    C [label="WIND"]

    D [label="NO", shape=plaintext]
    E [label="YES", shape=plaintext]
    F [label="YES", shape=plaintext]
    G [label="NO", shape=plaintext]

    A -> B [label="SUNNY"]
    A -> E [label="OVERCAST"]
    A -> C [label="RAIN"]

    B -> D [label="HIGH"]
    B -> F [label="NORMAL"]

    C -> F [label="WEAK"]
    C -> G [label="STRONG"]
}
"""

s = Source(dot_code, filename="decision_tree", format="png")
s.view()

```

Program 6

Build KNN Classification model for a given dataset

Screenshot

```
KNN:

import math
import matplotlib.pyplot as plt

def distance(p1, p2):
    return math.sqrt((p1[0] - p2[0])**2 + (p1[1] - p2[1])**2)

def knn(training_data, test_point, k):
    distances = []
    for point, label in training_data:
        d = distance(point, test_point)
        distances.append((d, label))
    distances.sort()
    k_nearest = distances[:k]
    labels = [label for _, label in k_nearest]
    prediction = max(set(labels), key=lambda label: label.count)
    return prediction

def visualization_knn(data, test_point, pred_label, new_point):
    colors = {'A': 'Blue', 'B': 'red'}
    markers = {'A': 'o', 'B': 's'}

    for point, label in data:
        plt.scatter(point[0], point[1], color=colors[label],
                    marker=markers[label], label=label)
    plt.scatter(test_point[0], test_point[1], color='green')
    if new_point is not None and new_label is not None:
        plt.scatter(new_point[0], new_point[1], color='orange')
```

```
plt.legend()
```

```
plt.grid(True)
```

```
plt.title('KNN class')
```

```
plt.xlabel('Feature 1')
```

```
plt.ylabel('Feature 2')
```

```
plt.show()
```

*2)

```
if name == "--main--"
```

```
data = [(1, 2, 'A'), (2, 3, 'A'), (3, 1, 'A'), (6, 5, 'B'),  
        (7, 7, 'B'), (8, 6, 'B')]
```

```
test = (2, 2)
```

```
result = knn(data, test, k=3)
```

```
print(result)
```

```
new_point = knn(data, new_point, k=3)
```

```
print(new_point)
```

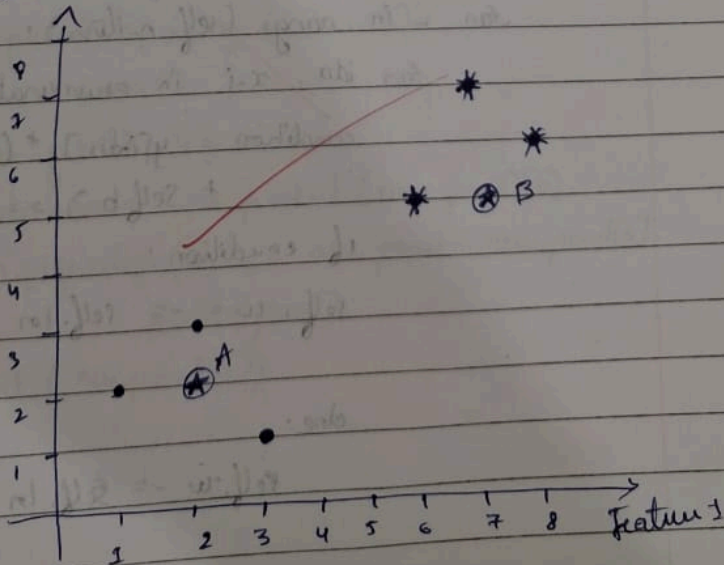
```
visualize_knn()
```

Obt:

Predicted class for test : A

Predicted class for new_point : B

Feature 2



code:

```
import math
import matplotlib.pyplot as plt

# Step 1: Distance calculation (Euclidean)
def distance(p1, p2):
    return math.sqrt((p1[0] - p2[0])**2 + (p1[1] - p2[1])**2)

# Step 2: KNN Function
def knn(training_data, test_point, k):
    distances = []
    for point, label in training_data:
        d = distance(point, test_point)
        distances.append((d, label))
    distances.sort()
    k_nearest = distances[:k]
    labels = [label for _, label in k_nearest]
    prediction = max(set(labels), key=labels.count)
    return prediction

# Step 3: Visualization Function
def visualize_knn(data, test_point, predicted_label, new_point=None,
new_label=None):
    colors = {'A': 'blue', 'B': 'red'}
    markers = {'A': 'o', 'B': 's'}

    # Plot training data
    for point, label in data:
        plt.scatter(point[0], point[1], color=colors[label],
marker=markers[label], label=label if f"train_{label}" not in
plt.gca().get_legend_handles_labels()[1] else "")

    # Plot test point
    plt.scatter(test_point[0], test_point[1], color='green', marker='*',
s=200, label=f'Test → {predicted_label}')

    # Plot new point if provided
    if new_point is not None and new_label is not None:
        plt.scatter(new_point[0], new_point[1], color='orange', marker='X',
s=150, label=f'New → {new_label}')
```

```
plt.legend()
plt.grid(True)
plt.title("KNN Classification")
plt.xlabel("Feature 1")
plt.ylabel("Feature 2")
plt.show()

# Step 4: Run everything
if __name__ == "__main__":
    # Training data
    data = [
        ([1, 2], 'A'),
        ([2, 3], 'A'),
        ([3, 1], 'A'),
        ([6, 5], 'B'),
        ([7, 7], 'B'),
        ([8, 6], 'B')
    ]

    # Test point
    test = [2, 2]
    result = knn(data, test, k=3)
    print("Predicted class for test:", result)

    # New point to classify
    new_point = [7, 5]
    new_result = knn(data, new_point, k=3)
    print("Predicted class for new point:", new_result)

    # Visualize
    visualize_knn(data, test, result, new_point=new_point,
new_label=new_result)
```


Program 7

Build Support vector machine model for a given dataset

Screenshot

```
SVM:

import numpy as np
import matplotlib.pyplot as plt

class SVM:
    def __init__(self, learning_rate=0.001, lambda_param=0.01, n_iters=1000):
        self.lr = learning_rate
        self.lambda = lambda_param
        self.n_iters = n_iters
        self.w = None
        self.b = None

    def fit(self, x, y):
        x = np.where(y <= 0, -1, 1)
        n_samples, n_features = x.shape
        self.w = np.zeros(n_features)
        self.b = 0

        for _ in range(self.n_iters):
            for idx, x_i in enumerate(x):
                condition = y[idx] * (np.dot(x_i, self.w) + self.b) >= 1
                if condition:
                    self.w -= self.lr * (2 * self.lambda * self.w + self.w)
                else:
                    self.w -= self.lr * (2 * self.lambda * self.w + self.w)
```

```

def predict(self, x):
    approx = np.dot(x, self.w) + self.b
    return np.sign(approx)

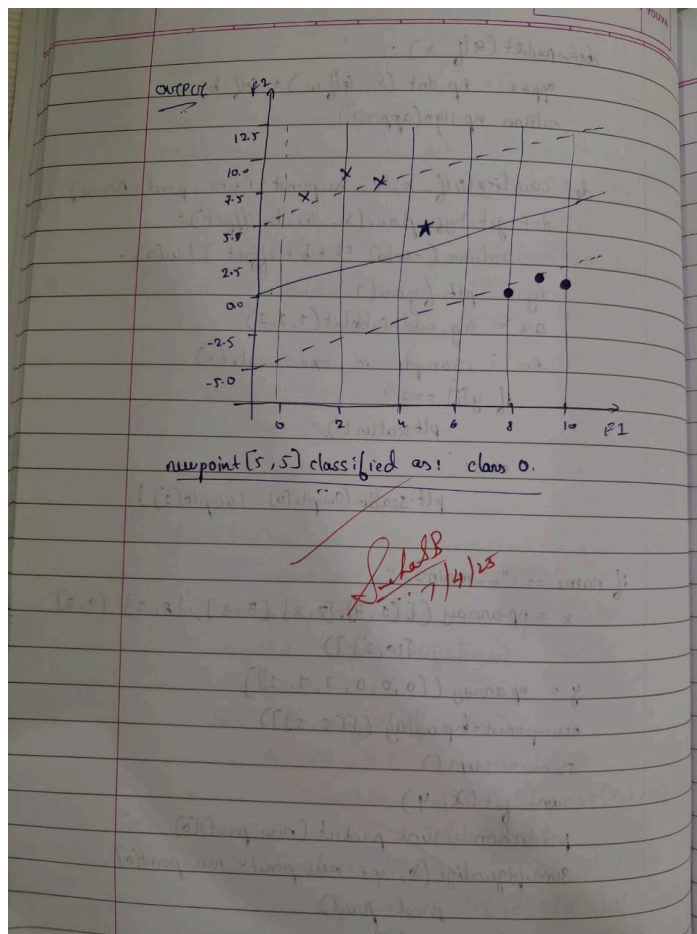
def visualize(self, x, y, newpoint=None, pred=None):
    def get_hyperspace(x, w, b, offset):
        return (-w[0] * x + b + offset) / w[1]

    fig = plt.figure()
    ax = fig.add_subplot(1, 1, 1)
    for i, sample in enumerate(x):
        if y[i] == 1:
            plt.scatter(i, sample)
        else:
            plt.scatter(i, sample)

    if newpoint is not None:
        plt.scatter(newpoint[0], newpoint[1])

    if name == "__main__":
        x = np.array([3, 3], [2, 8], [3, 8], [8, 2], [9, 2], [10, 2])
        y = np.array([0, 0, 0, 1, 1, 1])
        newpoint = np.array([5, 5])
        sum = sum()
        sum = f(x, y)
        prediction = sum.predict(newpoint[0])
        sum.visualize(x, y, newpoint=newpoint[0],
                     pred=pred)
        print(newpoint[0])

```



code:

```
import numpy as np
import matplotlib.pyplot as plt

class SVM:
    def __init__(self, learning_rate=0.001, lambda_param=0.01, n_iters=1000):
        self.lr = learning_rate
        self.lambda_param = lambda_param
        self.n_iters = n_iters
        self.w = None
        self.b = None

    def fit(self, X, y):
        y = np.where(y <= 0, -1, 1) # Convert labels to -1 and 1
        n_samples, n_features = X.shape
        self.w = np.zeros(n_features)
        self.b = 0

        for _ in range(self.n_iters):
            for idx, x_i in enumerate(X):
                condition = y[idx] * (np.dot(x_i, self.w) + self.b) >= 1
                if condition:
                    self.w -= self.lr * (2 * self.lambda_param * self.w)
                else:
                    self.w -= self.lr * (2 * self.lambda_param * self.w -
np.dot(x_i, y[idx]))
                    self.b += self.lr * y[idx]

    def predict(self, X):
        approx = np.dot(X, self.w) + self.b
        return np.sign(approx)

    def visualize(self, X, y, new_point=None, prediction=None):
        def get_hyperplane(x, w, b, offset):
            return (-w[0] * x + b + offset) / w[1]

        fig = plt.figure()
        ax = fig.add_subplot(1, 1, 1)

        # Plot existing data points
        for i, sample in enumerate(X):
```

```

        if y[i] == 1:
            plt.scatter(sample[0], sample[1], marker='o', color='blue',
label='Class +1' if i == 0 else "")
        else:
            plt.scatter(sample[0], sample[1], marker='x', color='red',
label='Class -1' if i == 0 else "")

    # Plot decision boundary
    x0 = np.linspace(np.min(X[:, 0])-1, np.max(X[:, 0])+1, 100)
    x1 = get_hyperplane(x0, self.w, self.b, 0)
    x1_m = get_hyperplane(x0, self.w, self.b, -1)
    x1_p = get_hyperplane(x0, self.w, self.b, 1)

    ax.plot(x0, x1, 'k-', label='Decision Boundary')
    ax.plot(x0, x1_m, 'k--', label='Margins')
    ax.plot(x0, x1_p, 'k--')

    # Plot the new point
    if new_point is not None:
        color = 'green' if prediction == 1 else 'orange'
        label = f'New Point: Class {"1" if prediction == 1 else "0"}'
        plt.scatter(new_point[0], new_point[1], c=color, s=100,
edgecolors='black', label=label, marker='*')

    ax.legend()
    plt.xlabel("Feature 1")
    plt.ylabel("Feature 2")
    plt.title("SVM with New Point Prediction")
    plt.grid(True)
    plt.show()

# 🚀 Example usage
if __name__ == "__main__":
    # Training data
    X = np.array([
        [1, 7],
        [2, 8],
        [3, 8],
        [8, 1],
        [9, 2],

```

```
[10, 2]
])
y = np.array([0, 0, 0, 1, 1, 1]) # 0 -> -1, 1 -> +1

# New point to classify
new_point = np.array([[5, 5]])

# Train and predict
svm = SVM()
svm.fit(X, y)
prediction = svm.predict(new_point)[0]

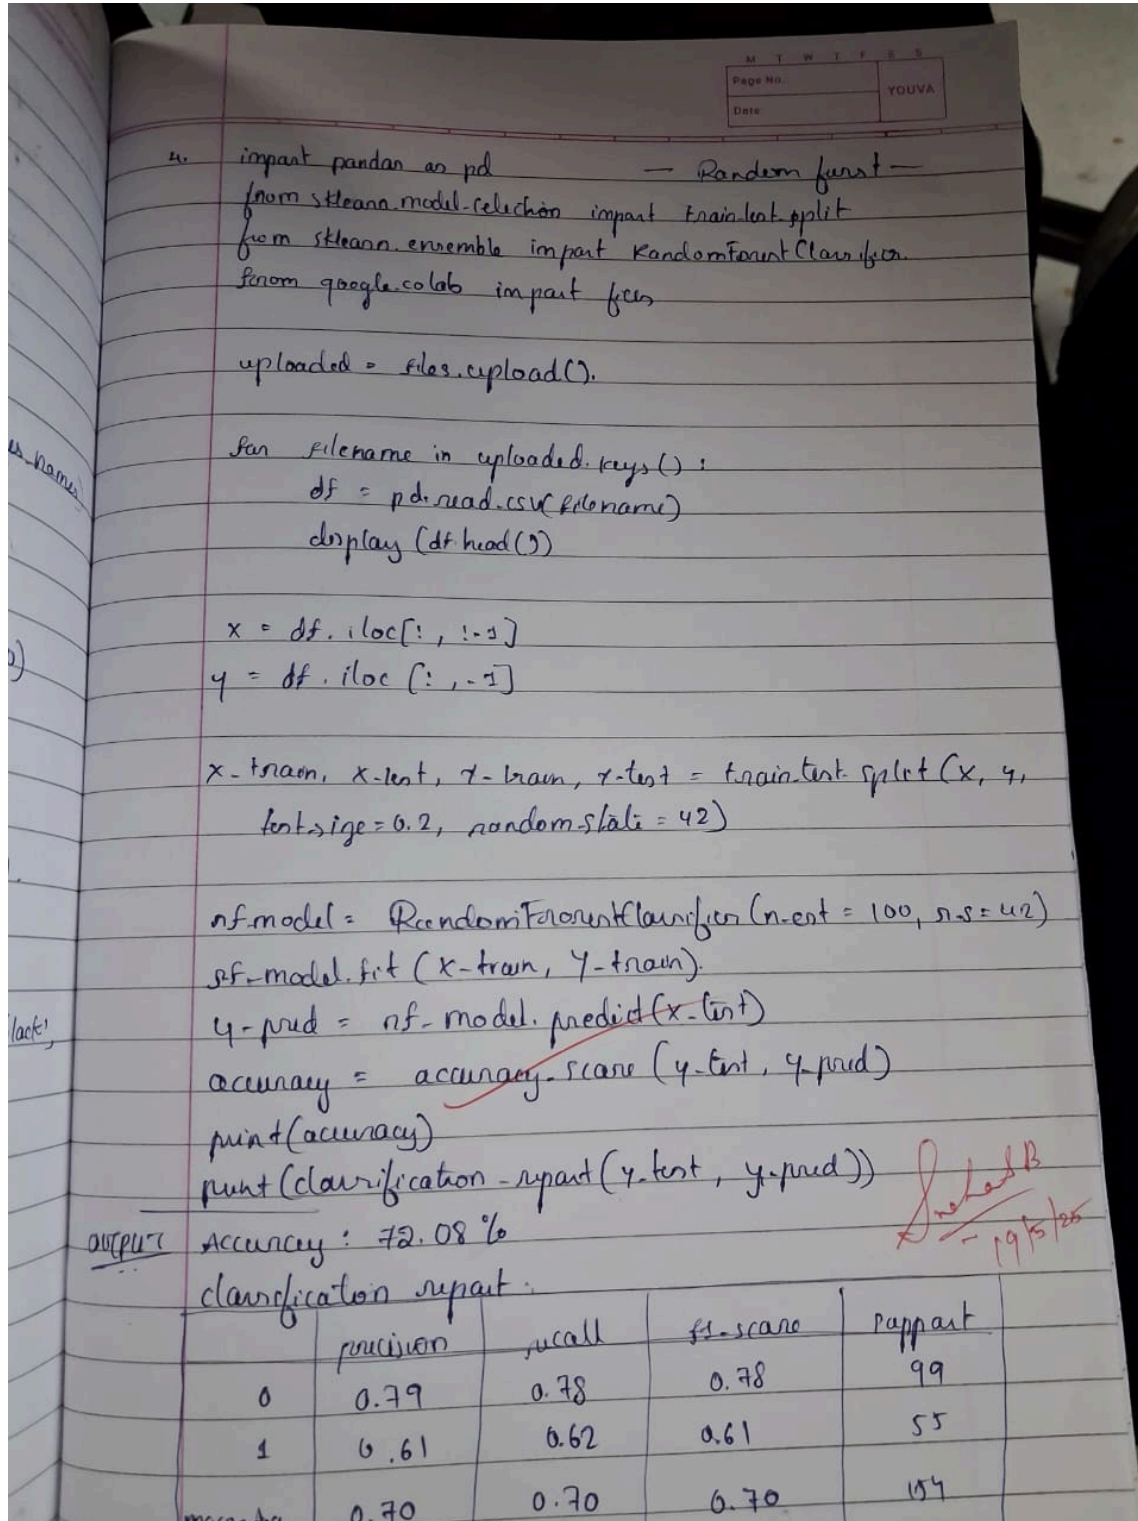
# Visualize
svm.visualize(X, y, new_point=new_point[0], prediction=prediction)

# Print prediction
print(f"New point {new_point[0]} classified as: {'Class 1' if prediction
== 1 else 'Class 0'}")
```

Program 8

Implement Random forest ensemble method on a given dataset

Screenshot



code:

```
#RANDOM FOREST
# STEP 1: Import required libraries
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, classification_report
from google.colab import files

# STEP 2: Upload your dataset
uploaded = files.upload()

# STEP 3: Load the dataset (assuming it's a CSV)
for filename in uploaded.keys():
    df = pd.read_csv(filename)
    print(f>Data loaded from: {filename}")
    display(df.head()) # Display first 5 rows of data

# STEP 4: Preprocessing
# Assume the last column is the target variable (label)
X = df.iloc[:, :-1] # Features (all rows, all columns except last)
y = df.iloc[:, -1]  # Target (last column)

# STEP 5: Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
random_state=42)

# STEP 6: Initialize and train the Random Forest model
rf_model = RandomForestClassifier(n_estimators=100, random_state=42) # 100
trees in the forest
rf_model.fit(X_train, y_train)

# STEP 7: Make predictions on the test set
y_pred = rf_model.predict(X_test)

# STEP 8: Evaluate the model
accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy of Random Forest Model: {accuracy * 100:.2f}%")

# STEP 9: Print classification report
print("Classification Report:")
print(classification_report(y_test, y_pred))
```

Program 9

Implement Boosting ensemble method on a given dataset

Screenshot

LAB-5

Decision Boost

```
1. import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.tree import DecisionTreeClassifier
from sklearn.datasets import make_classification
from sklearn.metrics import accuracy_score
from sklearn.decomposition import PCA
```

```
sns.set(style="whitegrid")
```

class Boost:

```
def __init__(self, n_est=50):
```

```
    self.n_est = n_est
```

```
    self.alphas = []
```

```
    self.models = []
```

```
    self.errors = []
```

```
def fit(self, X, y):
```

```
    n_samples, n_features = X.shape
```

```
    w = np.ones(n_samples) / n_samples
```

```
    for estimation in range(self.n_est):
```

```
        model = DecisionTreeClassifier(max_depth=1)
```

```
        model.fit(X, y, sample_weight=w)
```

```
        y_pred = model.predict(X)
```

```
        err = np.sum(w * (y_pred != y)) / np.sum(w)
```

```
        self.errors.append(err)
```

```
        alpha = 0.5 * np.log((1 - err) / err) if err < 1
        else 0
```

```
        self.alphas.append(alpha)
```


self.models.append(model)

$w = w + \text{np.exp}(-\alpha * y * y.\text{pred})$

$w = w / \text{np.sum}(w)$

def predict(self, x):

final_pred = np.zeros(x.shape[0])

for model, alpha in zip(self.models, self.alphas):

final_pred += alpha * model.predict(x)

return np.sign(final_pred)

def score(self, x, y):

return accuracy_score(y, self.predict(x))

x, y = make_classification(n_samples=500, n_features=2,
n_informative=2, n_redundant=2, n_classes=2)

$y = 2 * y - 1$

boost = Boost(n_estimators=50)

boost.fit(x, y)

accuracy = boost.score(x, y)

print f("model accuracy: {accuracy:.4f}")

output: accuracy: 0.9580

code:

```
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.tree import DecisionTreeClassifier
from sklearn.datasets import make_classification
from sklearn.metrics import accuracy_score
from sklearn.decomposition import PCA

# Set up plot style
sns.set(style="whitegrid")

class AdaBoost:
    def __init__(self, n_estimators=50):
        self.n_estimators = n_estimators
        self.alphas = [] # Weights of each weak classifier
        self.models = [] # Weak classifiers (e.g., decision stumps)
        self.errors = [] # List to store error for each estimator

    def fit(self, X, y):
        # Initialize weights for each data point
        n_samples, n_features = X.shape
        w = np.ones(n_samples) / n_samples # Equal weights initially

        for estimator in range(self.n_estimators):
            # Train weak classifier (decision stump)
            model = DecisionTreeClassifier(max_depth=1) # Decision stump
            model.fit(X, y, sample_weight=w)
            y_pred = model.predict(X)

            # Calculate error rate
            err = np.sum(w * (y_pred != y)) / np.sum(w)
            self.errors.append(err)

            # Compute alpha (weight for the classifier)
            alpha = 0.5 * np.log((1 - err) / err) if err < 1 else 0
            self.alphas.append(alpha)
            self.models.append(model)

            # Update weights for misclassified samples
            w = w * np.exp(-alpha * y * y_pred) # Update weights based on
```

```

classifier performance
    w = w / np.sum(w)  # Normalize the weights

    def predict(self, X):
        # Initialize the final prediction
        final_pred = np.zeros(X.shape[0])

        for model, alpha in zip(self.models, self.alphas):
            final_pred += alpha * model.predict(X)

        # Return the sign of the final prediction
        return np.sign(final_pred)

    def score(self, X, y):
        # Return accuracy of the model
        return accuracy_score(y, self.predict(X))

# Generate a synthetic binary classification dataset with 2 informative
features
X, y = make_classification(n_samples=500, n_features=2, n_informative=2,
n_redundant=0, n_classes=2, random_state=42)

# Convert labels to -1 and 1 for AdaBoost
y = 2 * y - 1

# Create and train AdaBoost model
adaboost = AdaBoost(n_estimators=50)
adaboost.fit(X, y)

# Evaluate the model
accuracy = adaboost.score(X, y)
print(f"Model accuracy: {accuracy:.4f}")

# Plot error over iterations
plt.figure(figsize=(10, 6))
plt.plot(range(1, adaboost.n_estimators + 1), adaboost.errors, marker='o',
linestyle='-', color='b')
plt.title('Error vs. Number of Estimators')
plt.xlabel('Number of Estimators')
plt.ylabel('Error')
plt.grid(True)

```

```
plt.show()

# Plot decision boundary for final model
x_min, x_max = X[:, 0].min() - 1, X[:, 0].max() + 1
y_min, y_max = X[:, 1].min() - 1, X[:, 1].max() + 1

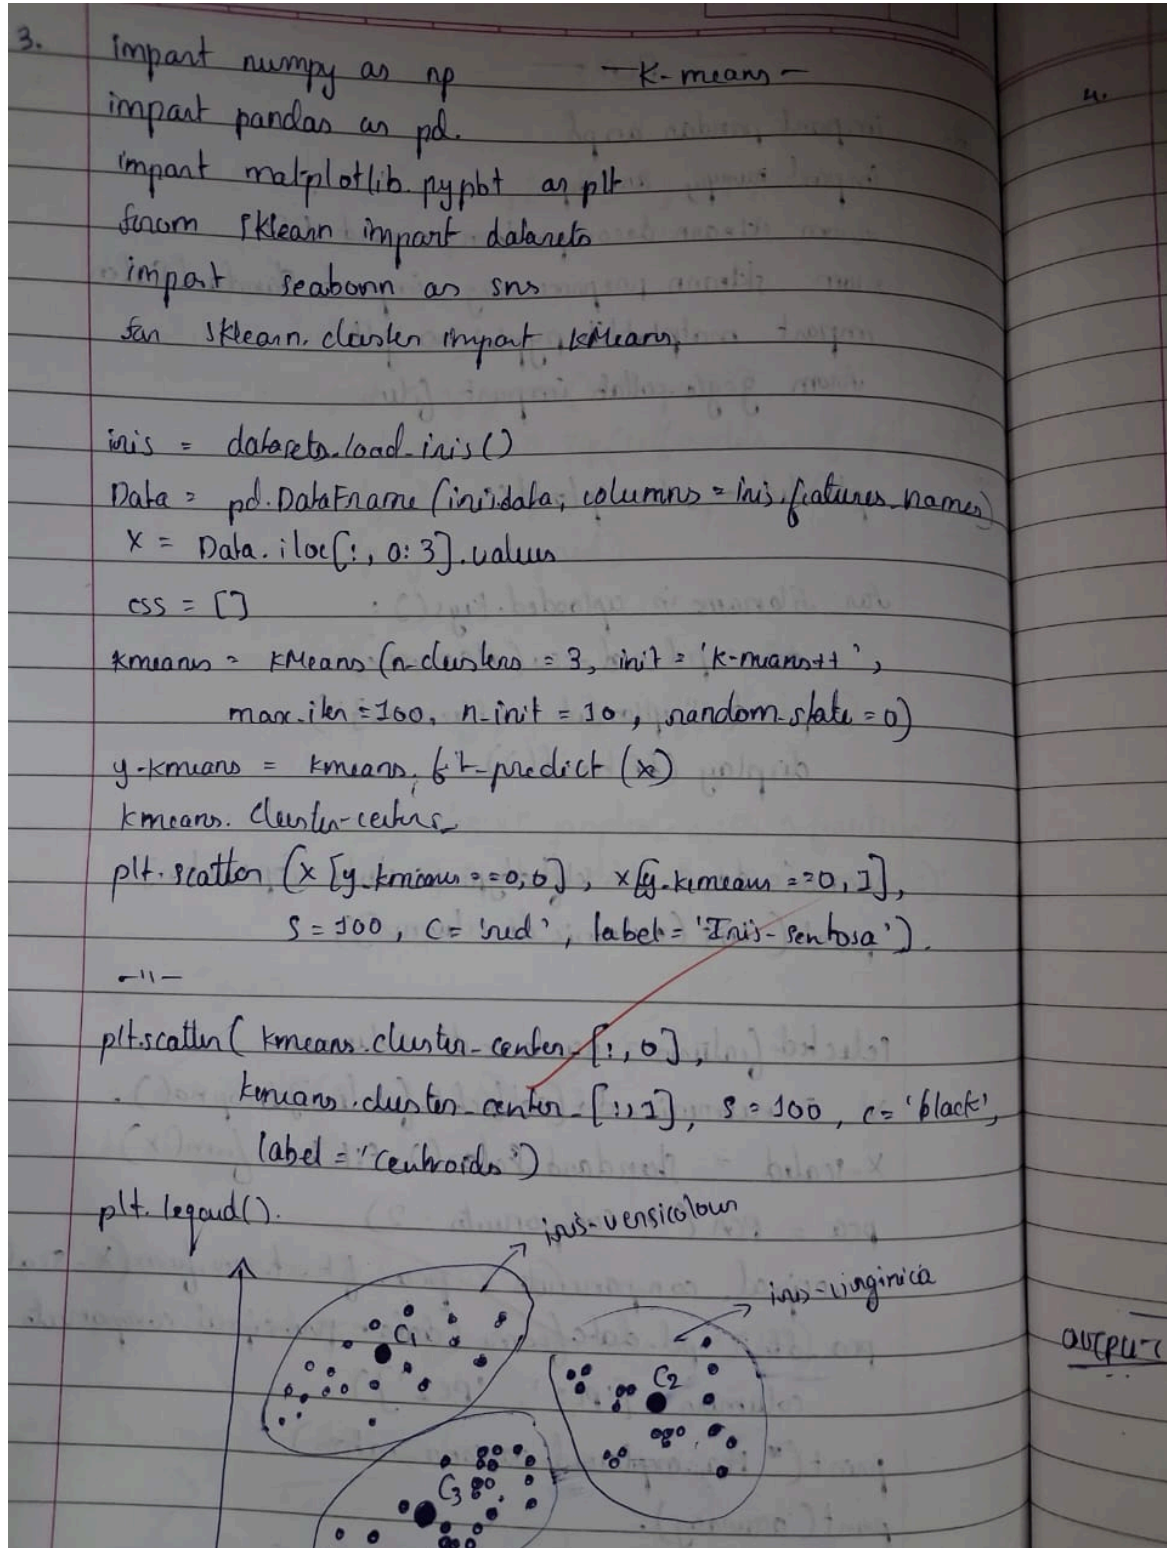
xx, yy = np.meshgrid(np.arange(x_min, x_max, 0.1),
                     np.arange(y_min, y_max, 0.1))
Z = adaboost.predict(np.c_[xx.ravel(), yy.ravel()])
Z = Z.reshape(xx.shape)

plt.figure(figsize=(10, 6))
plt.contourf(xx, yy, Z, alpha=0.75, cmap='coolwarm')
plt.scatter(X[:, 0], X[:, 1], c=y, edgecolors='k', marker='o', s=50,
            cmap='coolwarm')
plt.title('AdaBoost Decision Boundary')
plt.xlabel('Feature 1')
plt.ylabel('Feature 2')
plt.show()
```

Program 10

Build k-Means algorithm to cluster a set of data stored in a .CSV file

Screenshot



code:

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn import datasets
import seaborn as sns
from sklearn.cluster import KMeans

iris = datasets.load_iris()
print("Dataset loaded successfully")

Data = pd.DataFrame(iris.data, columns = iris.feature_names)

#Top values of Dataset
# Data.head()
x=Data.iloc[:,0:3].values

css=[]

# # Finding inertia on various k values
# for i in range(1,8):
#     kmeans=KMeans(n_clusters = i, init = 'k-means++',
#                   max_iter = 100, n_init = 10, random_state = 0).fit(x)
#     css.append(kmeans.inertia_)
#Applying Kmeans classifier
kmeans = KMeans(n_clusters=3,init = 'k-means++', max_iter = 100, n_init = 10,
random_state = 0)
y_kmeans = kmeans.fit_predict(x)
kmeans.cluster_centers_

# Visualising the clusters - On the first two columns
plt.scatter(x[y_kmeans == 0, 0], x[y_kmeans == 0, 1],
            s = 100, c = 'red', label = 'Iris-setosa')
plt.scatter(x[y_kmeans == 1, 0], x[y_kmeans == 1, 1],
            s = 100, c = 'blue', label = 'Iris-versicolour')
plt.scatter(x[y_kmeans == 2, 0], x[y_kmeans == 2, 1],
            s = 100, c = 'green', label = 'Iris-virginica')

# Plotting the centroids of the clusters
plt.scatter(kmeans.cluster_centers_[:, 0], kmeans.cluster_centers_[:,1],
            s = 100, c = 'black', label = 'Centroids')
plt.legend()
```

Program 11

Implement Dimensionality reduction using Principal Component Analysis (PCA) method

Screenshot

M T W T F S S
Page No. YOUVA
Date:

PCA

```

2. import pandas as pd
import numpy as np
from sklearn.decomposition import PCA
from sklearn.preprocessing import StandardScaler
import matplotlib.pyplot as plt
from google.colab import files

uploaded = files.upload()

for filename in uploaded.keys():
    df = pd.read_csv(filename)
    print(f"uploaded: {filename}")
    display(df.head())

    numeric_df = df.select_dtypes(include=[np.number])
    print(list(numeric_df.columns))

    selected_features = numeric_df.columns
    X = numeric_df[selected_features].dropna()
    X_scaled = StandardScaler().fit_transform(X)
    pca = PCA(n_components=2)
    principal_components = pca.fit_transform(X_scaled)
    pca_df = pd.DataFrame(data=principal_components,
                          columns=['PC1', 'PC2'])
    print("PCA explained variance ratio")
    print(accuracy)

    output: variance_ratio: [0.52163044 0.28631263]
    Model accuracy: 0.9580
  
```

code:

```

# STEP 1: Import packages
import pandas as pd
import numpy as np
from sklearn.decomposition import PCA
from sklearn.preprocessing import StandardScaler
import matplotlib.pyplot as plt
from google.colab import files

# STEP 2: Upload the CSV file
uploaded = files.upload()

# STEP 3: Load the dataset
for filename in uploaded.keys():
    df = pd.read_csv(filename)
    print(f"✅ Uploaded: {filename}")
    display(df.head())
  
```

```

# STEP 4: Select numerical columns
numeric_df = df.select_dtypes(include=[np.number])
print("🇩🇪 Numerical features found:", list(numeric_df.columns))

# OPTIONAL: Manually select columns if needed
# selected_features = ['feature1', 'feature2', ...]
selected_features = numeric_df.columns # use all numeric features for now

# STEP 5: Standardize data
X = numeric_df[selected_features].dropna()
X_scaled = StandardScaler().fit_transform(X)

# STEP 6: Apply PCA
pca = PCA(n_components=2)
principal_components = pca.fit_transform(X_scaled)

# STEP 7: Create DataFrame for components
pca_df = pd.DataFrame(data=principal_components, columns=['PC1', 'PC2'])

# STEP 8: Visualize the first two principal components
plt.figure(figsize=(8,6))
plt.scatter(pca_df['PC1'], pca_df['PC2'], alpha=0.7)
plt.xlabel('Principal Component 1')
plt.ylabel('Principal Component 2')
plt.title('2D PCA Visualization')
plt.grid(True)
plt.show()

# STEP 9: Explained variance ratio
print("✅ Explained Variance Ratio:", pca.explained_variance_ratio_)

print(f"Model accuracy: {accuracy:.4f}")

```